

Request


Hermes

### Coinference-aware Request Annotator (§4.2)

Offline Profiling → Knowledge Base  Online Analysis

Create new Coinference or add request to existing coinference

### Inter-Coinference Scheduler (§4.3)



Find the next request to serve

### Intra-Coinference Resource Provisioner (§4.4)


Clairvoyant Cache Management (§4.4.1)

Clairvoyant Queuing on Non-LLM Backends (§4.4.2)


Clairvoyant Backend Selection (§4.4.3)

### LLM Backends


LLM Backend 1



LLM Backend 2




CPU Scheduler



Network Scheduler



GPU Scheduler



Non-LLM Backends