

Request

Hermes

Coinference-aware Request Annotator (§4.2)

Offline Profiling


Knowledge Base 

Online Analysis

Create new Coinference or add request to existing Coinference

Inter-Coinference Scheduler (§4.3)

Coinference Queue



Find the next request to serve

Intra-Coinference Resource Provisioner (§4.4)

Clairvoyant Cache Management (§4.4.1)

Clairvoyant Queuing on Non-LLM Backends (§4.4.2)

Clairvoyant Backend Selection (§4.4.3)

LLM Backend 1

LLM

KV Cache

LoRA

LLM Backends

CPU

Scheduler



Network

Scheduler



GPU

Scheduler



Non-LLM Backends